

Quantum Computation

quinten tupker

October 8 2020 - January 12, 2021

Introduction

These notes are based on the course lectured by Professor Richard Jozsa in Michaelmas 2020. Due to the measures taken in the UK to limit the spread of Covid-19, these lectures were delivered online. These are not meant to be an accurate representation of what was lectures. They solely represent a mix of what I thought was the most important part of the course, mixed in with many (many) personal remarks, comments and digressions... Of course, any corrections/comments are appreciated.

This course is meant to be a second course on quantum computation. In particular, all the prerequisite knowledge is covered in Cambridge's Part II Quantum Information and Computation course. Lecture notes for this course can be found online.

Now, to start describing course content. Quantum Computation studies algorithms that can be run on quantum computers. Although they have yet to be implemented in practice (but are definitely being developed at a rapid pace), and in particular how they differ with classical computation. A remarkable result, is that at least superficially, quantum algorithms appear to be more powerful than classical algorithms, although it remains unclear if this is definite fact, or if it simply easier for humans to solve complex problems using quantum algorithms instead of classical algorithms. This is remarkable in many ways in a philosophical sense, but here we focus on how to take advantage of these changes. As such, we will begin with a review and extension of one of the most famous quantum algorithms, which is Schur's factoring algorithm.

1 Schur's Algorithm Revisited and the Hidden Hidden Subgroup Problem

Schur's factoring algorithm finds a factor for an arbitrary number N . It does not perform a complete factorisation. It merely computes a factor, which of course can be repeated arbitrarily to get a complete factorisation, but that is not the point here. The complexity of such an algorithm is typically measured in terms of the number of digits of N , which we may denote $n = \ln(N)$. In these

terms Schur's algorithm is $O(n^3)$, which means it is "efficient" (computationally feasible) or

Definition 1 (efficient algorithm). An algorithm is efficient if it runs in polynomial time, which generally means it is considered doable in practice.

By comparison, the fastest known classical algorithm runs in $O(e^{n^{1/3} \ln(n)^{1/3}})$. Anyways, here is an outline of Schur's algorithm:

1. choose $a < N$ such that $(a, N) = 1$ (coprime). This can be done efficiently, since the probability of a being coprime is fixed, and we can quickly calculate the GCD using Euclid's algorithm. Then consider $f(x) = a^x \pmod{N}$.
2. Use quantum algorithms to calculate the period of the this function (so we have converted a factoring problem into period finding). Since a is coprime, it is guaranteed to be periodic.
3. compute the factor using number theory

The crucial component here is the period finding, which cannot be done efficiently using a classical algorithm. So let's review quantum period finding. Also, note that as usual, quantum oracles are implemented as unitary operators by converting $f : \mathbb{Z}_M \rightarrow \mathbb{Z}_N$ to $U_f |x\rangle |0\rangle \mapsto |x\rangle |f(x)\rangle$. Then, if f has period r (unknown), and f is one-to-one on every period, then period finding can be done using

1. make $\frac{1}{\sqrt{M}} \sum_0^{M-1} |i\rangle |0\rangle$
2. apply U_f
3. measure the output register to get

$$\frac{1}{\sqrt{A}}(|x_0\rangle + |x_0 + r\rangle + \cdots + |x_0 + (A-1)r\rangle) |f(x_0)\rangle$$

Now the next step is the tricky part, and really is what uses the "quantum magic" here, and that involves the use of the Quantum Fourier Transform (QFT). [This ends lecture 1]

4. We then apply the QFT, which maps $|k\rangle \mapsto \sum_{y=0}^{M-1} e^{xy} |y\rangle$, which, after some calculation (use $\sum e^{2\pi kx/y} = y\delta_{xy}$) leaves us with

$$\text{QFT}|\text{per}\rangle = \sqrt{A/M} \sum_{k=0}^{r-1} \omega^{x_0 k M/r} |kM/r\rangle$$

5. Making a measurement we get $C = k_0 M/r$, so $\frac{k_0}{r} = \frac{C}{M}$. If k_0, r are coprime, we are done, since we can reduce $\frac{C}{M}$ to simplest terms (use Euclid's algorithm to cancel out the gcd). Now, number theory tells

us that the probability of being coprime is finite and shrinks slowly (as $O(1/\log \log(M))$), and so we can just repeat until we get the right period. Since f is one-to-one on each period, it is easy to check if our period is correct.

I feel that just being able to check if the period is correct is a somewhat lame reason to require that the function be one-to-one on each period, but improvements although not difficult, would complicate this explanation.

Anyways, let's see if we can motivate the Quantum Fourier Transform a bit better. The challenge we face is that our state, $|R\rangle$ takes the form

$$|R\rangle = \sum_k a_k |x_0 + kr\rangle$$

for an arbitrary x_0 . In other words, we have an arbitrary shift that we want to ignore somehow. How do we do that? A natural way to spot "things that ignore shifts" would be to define the shift operator

$$U|x\rangle = |x + 1 \bmod M\rangle$$

and then, how do we say, "we don't care about U "? We look for the eigenvectors of U , which are by definition, the states least affected by U . Fortunately, U is a permutation matrix, so unitary, so is a quantum gate. Then, the eigenbasis of U is what we may call the set of shift-invariant states χ_k . If we write R in terms of this basis we are bound to get a state of the form

$$\sum_k a_k \lambda_k^{x_0} |\chi_k\rangle$$

where $\mathbb{P}(k) = |a_k \lambda_k^{x_0}|^2 = |a_k|^2$ since as eigenvalues of a unitary matrix, $|\lambda_k| = 1$ always. So as expected, probabilities are preserved (this is not that crucial - but it's important they don't differ that much). Anyways, important is that this transformation allows us to express our state as a sum of multiples of the period, which is what we want.

All that remains is to find this basis, and the operation that expresses them in terms of it. As eigenvectors of a unitary matrix are orthogonal, all we need to do is zip the eigenvectors into a matrix. These eigenvectors are just of the form $e^{2\pi i k l / M}$, so we get the Quantum Fourier Transform we expect. [End of lecture 2]

1.1 The Hidden Subgroup Problem (HSP)

The hidden subgroup problem asks the question how can we find subgroup K of group G (this course only considers finite G) given a function $f : G \rightarrow X$ that is an invertible function of the left cosets of K . Our goal is to solve this problem in $O(\text{poly}(\ln(|G|)))$.

Examples of this include Schur's algorithm where we have $G = \mathbb{Z}_p^*$, then $K = 0, r, 2r, \dots$, and then f used before works for our purpose. Another example is calculating **discrete logarithm**, which if done efficiently could break

encryption methods. This involves calculating logarithms on \mathbb{Z}_p^* , so given x finding y such that for group generator g , $x = g^y$. To formulate this as an HSP consider

$$\begin{aligned} f : \mathbb{Z}_{p-1} \times \mathbb{Z}_{p-1} &\rightarrow \mathbb{Z}_p^* \\ (a, b) &\mapsto g^a x^{-b} = g^{a-yb} \end{aligned}$$

where we can see that $f(a_1, b_1) = f(a_2, b_2)$ iff $(a_2, b_2) - (a_1, b_1) = \lambda(y, 1)$ some λ . So using $K = \{\lambda(y, 1) | \lambda \in \mathbb{Z}_{p-1}\}$ works.

Those both involve abelian groups, but a big area of interest is solving the problem for non-abelian groups. Examples of such problems include finding the Automorphism group of a graph. For graph A this means finding $\text{Aut}(A)$ a subgroup of the group of permutations of the vertices of A such that the overall structure (edges) are preserved. This can be formulated quite naturally into an HSP by taking X to be the set of all n vertex graphs, and then to consider the function $f_A(\pi)$ which applies permutation π to A . Clearly the result depends only on the coset of $\text{Aut}(A)$.

Even more famous is the **Graph Isomorphism** problem (GI), which is to check whether or not two labelled graphs are isomorphic (so whether there exists a permutation turning one into the other). This can be converted into an HSP, although the process is more complicated. There also exists few good classical algorithms, although in 2017 someone found an algorithm doing it in quasi-polynomial time.

But how far will we get with quantum algorithms. The abelian case has been fully solved, but unfortunately the non-abelian case (which the above two problems belong to) remains an important unsolved problem in the field. [End of lecture 3]

Briefly, another example of an HSP involving the Dihedral group is how to calculate the shortest vector within a lattice. No good quantum algorithm exists to solve this problem.

Now, let's develop the formalism to describe the quantum algorithm to solve the hidden subgroup problem for any finite abelian group. Contrary to my expectations, we jump straight into representation theory. Well, that is not really contrary to my expectations, but I had imagine it would be more natural to look at the problem as permutations, and to define invariant states from there. Apparently, while it is not mentioned, it does not seem to be the central approach thought of here.

Instead we rely on the fact that for a finite abelian group, all representations can be seen as one dimensional representations $\chi : G \rightarrow S^1$, the unit circle embedded in the complex plain. Being one dimensional they are all irreducible, and also satisfy the following three properties:

- $\chi(g) = e^{2\pi i k / |G|}$ some $k \in \mathbb{Z}$
- $\langle \chi_i, \chi_j \rangle = \frac{1}{|G|} \sum_{g \in G} \chi_i(g) \chi_j^*(g) = \delta_{ij}$

- There are exactly $|G|$ distinct such representations, which we may label χ_g
- the restriction of an irreducible representation to a subgroup is also irreducible

A particular application of this is that **trivial representation** $\chi_0(g) = 1$ is orthogonal to all other representations, so in general $\sum_g \chi_k(g) = 0$. This will be helpful later on. Also, note that since this is an abelian group, we will use $+$ to denote the group operation.

Now the crucial advantage that the abelian case has over the non-abelian case is that all elements commute, meaning in particular that all representations can be simultaneously diagonalised, which means a universal shift invariant state exists. In particular, we notice that

$$|\chi_k\rangle = \frac{1}{\sqrt{|G|}} \sum_{g \in G} \overline{\chi_k(g)} |g\rangle$$

Here the complex conjugate is really just convention, but it makes things look similar to Schur's algorithm. Here it is easy to see that the eigenvalues of a shift by g , denoted $U(g)$ is $U_k(g) |\chi_k\rangle = \chi_k(g) |\chi_k\rangle$.

The Quantum Fourier Transform (QFT) then as expected swaps the standard unit basis with the shift invariant basis meaning that $QFT |\chi_g\rangle = |g\rangle$ so considering components appropriately, and transposing to get the inverse we find

$$QFT |g\rangle = \frac{1}{\sqrt{|G|}} \sum_{k \in G} \chi_k(g) |k\rangle$$

Now, since we can classify all finite abelian groups as the product of cyclic groups, the following result on representations of such groups is quite valuable

Example 1. For $G = \mathbb{Z}_M$, $\chi_a(b) = e^{2\pi i ab/M}$ describes all representations of G . Even better is that for $G = \mathbb{Z}_{m_1} \times \cdots \times \mathbb{Z}_{M_r}$, $\chi_a(b) = e^{2\pi i(a_1 b_1/M_1 + \cdots + a_r b_r/M_r)}$. Furthermore, in this case

$$QFT_G = QFT_{M_1} \otimes \cdots \otimes QFT_{M_r}$$

Now let's start to truly describe the HSP algorithm for finite abelian groups. Here we are given oracle $f : G \rightarrow K$ and we work on state space $\mathcal{H}_{|G|} \otimes \mathcal{H}_{|K|}$

1. form the state space $\frac{1}{\sqrt{|G|}} \sum_{g \in G} |g\rangle |0\rangle$
2. Apply U_f for form $\frac{1}{\sqrt{|G|}} \sum_{g \in G} |g\rangle |f(g)\rangle$
3. measure second register to get $|g_0 + K\rangle$ where $|K\rangle = \frac{1}{\sqrt{|K|}} \sum_{k \in K} |k\rangle$.

4. Apply QFT and measure to get some state χ_g . Note that here g does not depend on g_0 due to the QFT, which is the strength of this algorithm. [End of lecture 4] From here we can identify representations that become trivial when restricted to K since we see that

$$\begin{aligned} QFT |K\rangle &= \frac{1}{\sqrt{|G||K|}} \sum_{l \in G} \left(\sum_{k \in K} \chi_l(k) \right) |l\rangle \\ &= \sqrt{\frac{|K|}{|G|}} \sum_{l \in G \text{ st } \chi_l|_K \text{ is trivial}} |l\rangle \end{aligned}$$

using the fact that $\sum \chi_l(k)$ is either $|K|$ or 0 depending on whether or not χ_l becomes trivial.

That ends the quantum part of the algorithm, where one essentially identifies the representation χ_g that is irreducible on G but that becomes the trivial representation when restricted to K . This is enough to find the subgroup of interest.

Example 2. For example, one can show that if K has generators k_1, \dots, k_M for $K = O(\ln(|K|)) = O(\ln(|G|))$ then if we repeat the algorithm $O(\ln(|G|))$ times we can find the generators (so the subgroup) with probability $> 2/3$.

Example 3. In the $G = \mathbb{Z}_{M_1} \times \dots \times \mathbb{Z}_{M_r}$, one can find that finding the group is equivalent to solving a set of linear equations.

Now what goes wrong in the non-abelian case? Essentially two things can break down. Firstly, we might not be able to implement the quantum fourier transform in an efficient way. Secondly, because elements don't commute, we cannot simultaneously (block) diagonalise the basis into a set of shift-invariant states. This means one cannot fully recover the representations as one might wish. It has nevertheless been proven when the subgroup is normal, one can solve this efficiently (Hallgren, Russel, Ta-Shma 2003 SIAM J Comp The Hidden Subgroup Problem and Quantum Computation Using Group Representations), and that in fact, even in the most general non-abelian case the information we deduce is sufficient to find K , however there is no efficient way to deduce K from the information we can gain using the more coarse-grained approach here (Ettinger, Hoyer and Knill, Hidden Subgroup States are Almost Orthogonal) (so query complexity is not high, but deducing the right information from it is high). [End of lecture 5]

2 Quantum Phase Estimation

The next problem/aglorithm we discuss is **Quantum Phase Estimation** which is how to find an eigenvalue $e^{2\pi i \phi}$ of eigenstate $|v_\phi\rangle$ of an arbitrary unitary operator U . To do so we need the controlled operator $c - U$ such that $c - U |0\rangle |\xi\rangle = |0\rangle |\xi\rangle$ and $c - U |1\rangle |\xi\rangle = |1\rangle U |\xi\rangle$.

Now the issue is that we cannot actually construct $c - U$ given U if we only have one of its eigenstates. If we have a full implementation, for example in the form of a circuit, we could simply replace every gate with a controlled gate, and that would do the job, however, if we just have a black box operator, there is an ambiguity that arises from the fact that we may consider U and $e^\alpha U$ to be same operator. If so we find that

$$c - e^{i\alpha}U(|0\rangle + |1\rangle)|\xi\rangle = |0\rangle|\xi\rangle + |1\rangle e^{i\alpha}U|x\rangle$$

so these differ by a relative phase, meaning the operator $c - U$ is truly different from $c - e^{i\alpha}U$ despite U and $e^{i\alpha}U$ being the “same”. To resolve this, it turns out it is sufficient as long as we have one extra eigenstate $|\alpha\rangle$ with eigenvalue $e^{i\alpha}$, since then we can implement the algorithm as follows:




Figure 1: Controlled U

Here the circle with the arrows indicate controlled swap operations between $|\alpha\rangle$ and $|\xi\rangle$, double lines indicate an arbitrary k -qubit state, while a single line denotes a single qubit state. X denotes the swap operator between $|0\rangle$ and $|1\rangle$ and $f(\alpha)$ denotes the matrix

$$\begin{pmatrix} 1 & 0 \\ 0 & e^{i\alpha} \end{pmatrix}.$$

Considering case-by-case we then find that

- when the control is $|0\rangle$ we get $e^{i\alpha - i\alpha} |0\rangle |\xi\rangle$.
- when the control is $|1\rangle$ we get $|1\rangle |\xi\rangle$

(typo in the picture, θ should be $-\alpha$). Now for our actual algorithm we don't need $c - U$ instead we need (what I'll call) $N - U$ which takes the form $N - U |x\rangle |\xi\rangle = |x\rangle U^x |\xi\rangle$. This can be implemented as seen below for decimal expansion $x = x_{n-1} \dots x_1 x_0$.




Figure 2: Generalised Controlled U

Now that we've implemented that, the actual quantum phase estimation algorithm is quite easy to implement:

1. start with $|+\rangle |v_\phi\rangle$ for $|+\rangle = H^n |0\rangle$
2. apply $N - X$ to get $\frac{1}{2^{n/2}} \sum_x e^{2\pi i \phi x} |x\rangle = |A\rangle$.
3. Apply $QFT_{2^n}^{-1}$ and measure to get $y_{n-1} \dots y_1 y_0$. Our estimate for ϕ the is

$$\phi \approx 0.y_0 \dots y_{n-1}$$

This is exact if $\phi = k/2^n$ some $0 \leq k < 2^n, k \in \mathbb{Z}$. If it's not, we have the following handy result

Theorem 1. If the algorithm yields estimate $\theta \approx \phi$ then

- the probability of θ being the estimate to ϕ that is closest possible to ϕ in binary is at least $4/\pi^2 \approx 0.4$.
- $\mathbb{P}(|\phi - \theta| \geq \epsilon) \leq \frac{1}{2^{n+1}\epsilon}$.

In particular, we notice that if we want at least m bit accuracy with certainty $1 - \eta$ then we should run the algorithm with n bits such that $n = m + \log_2(1/\eta)$. It is interesting here, to me, that the quantity n is sum of the desired accuracy m and the chance of that accuracy (expressed in binary) $\log_2(1/\eta)$. The fact these are equivalent in a sense is curious. [End of lecture 6]

Proof. First note that

$$QFT^{-1} |A\rangle = \frac{1}{2^n} \sum_y \sum_x e^{2\pi i (\phi - y/2^n)x} |y\rangle$$

So in particular, the probability of observing the value y is

$$\mathbb{P}(y) = \frac{1}{y^{2n}} \left| \sum_x e^{2\pi i (\phi - y/2^n)x} \right|^2 = \frac{1}{2^{2n}} \left| \frac{1 - e^{2\pi i \delta(y)}}{1 - e^{2\pi i \delta(y)}}$$

where $\delta(y) = \phi - y/2^n$. We use the following estimates to establish bounds on this quantity. Firstly, $|1 - e^{i\alpha}| = 2|\sin(\alpha/2)| \geq \frac{2}{\pi}|\alpha|$ if $|\alpha| < \pi$ and $|1 - e^{i\beta}| \leq \beta$ which can both be seen geometrically.

Then for our first estimate, by bounding appropriately we see that

$$\mathbb{P}(y) \geq \frac{1}{2^n} \left(\frac{2^{n+1}\delta(y)}{2\pi\delta(y)} \right)^2 = \frac{4}{\pi^2}$$

as required.

Our other estimate is a bit trickier. We see that

$$\mathbb{P}(y) \leq \frac{1}{2^{2n}} \left(\frac{2}{4\delta(y)} \right)^2 \leq \frac{1}{2^{2n+2}\delta(y)^2}$$

But we now also have to sum over all y such that $|\delta(y)| > \epsilon$. To do so we see that in general $\delta(y)$ takes the form $\delta(y) = \delta_{\pm} + k/2^n$ meaning that $|\delta(y)| \geq \epsilon + k/2^n$, so

$$\mathbb{P}(|\delta(\gamma)| \geq \epsilon) \leq 2 \sum_{k=0}^{\infty} \frac{1}{2^{2n+2}(\epsilon + k/2^n)^2} \leq \frac{1}{2} \int_0^{\infty} \frac{1}{(2^n\epsilon + k)^2} dk = \frac{1}{2^{n+2}\epsilon}$$

as required. \square

Now a few remarks. Firstly, if $c - U^{2^k}$ is implemented as $(c - U)^{2^k}$ then this algorithm is exponentially slow, however, for some special U , such as simple exponentiation $a \mapsto a^p$ this can be done in polynomial time. Secondly, if we apply the unitary part U_{PE} to an arbitrary state, we get a superposition in terms of eigenstates, but when we make a measurement, the estimate for every eigenstate is good. It is just that it is random which eigenstate we are actually measuring.

This second perspective is handy, as we can now express getting m digits correct with probability $1 - \eta$ as

$$U_{PE} |\xi\rangle = \sqrt{1 - \eta} |\phi_j\rangle + \sqrt{\eta} |\phi_{\perp}\rangle$$

where ϕ_j are all states with m correct digits, and ϕ_{\perp} are the incorrect states.

Finally, we note that Dr Jorza, in his notes, also describes how one might implement a quantum fourier transform in an arbitrary dimension (not necessarily a power of two) using quantum phase finding.

3 Amplitude Amplification

Our next algorithm we look at is a generalisation of the essence of Grover's search algorithm. For such, we start with some notation. We use P_A to denote projection onto the linear space A which if spanned by orthogonal basis α_i can be calculated as $P_A = \sum_i |\alpha_i\rangle \langle \alpha_i|$. We can then write a reflection as $I_A = I - P_A$. Finally, we note for convenience that for any unitary U , $UI_AU^{\dagger} = U_{UA}$ [End of lecture 7].

3.1 Grover's Algorithm

We begin by reviewing Grover's algorithm (which can be found in more detail in the part II notes of the course). Here we are given a function

$$f(x) = \begin{cases} 1 & x = x_0 \\ 0 & \text{else} \end{cases}$$

and are trying to find x_0 . This is an unstructured search really, and closely relate to NP complexity in classical complexity theory via the Boolean satisfiability problem. Now, to run this algorithm we define the **Grover rotation** $Q = -H_n I_{|0\rangle} H_n I_{|x_0\rangle} = -I_{|\psi_0\rangle} I_{|x_0\rangle}$. The result we rely on, but we do not prove here in particular is that

Theorem 2. Q is a rotation on the space spanned by $|\psi_0\rangle, |x_0\rangle$ by angle 2α where $\sin(\alpha) = \frac{1}{\sqrt{N}}$

We then proceed to run the algorithm via

1. Form the uniform state $|\psi_0\rangle$ (using Hadamard gates)
2. Apply Q m times where

$$m = \frac{\arccos(1/N)}{2 \arcsin(1/N)} = \frac{\beta}{2\alpha} \approx \frac{\pi}{4} \sqrt{N}$$

where β is the angle between $|\psi_0\rangle$ and $|x_0\rangle$ and the approximation applies when N is big.

3. Measure the state to get $|x_0\rangle$ with high probability $\cos^2(\alpha) \approx 1 - \frac{1}{N}$.

As we can see, this runs in $O(\sqrt{N})$ queries to the oracle. It also does something that a classical computer is not capable of, since in particular for $N = 4$ this is deterministic and requires exactly one query, which cannot be matched classically.

3.2 Returning to Amplitude Amplification

Amplitude amplification is a generalisation of Grover's algorithm. In particular, if we use G to denote the "good" subspace of \mathcal{H} then see that $\forall |\psi\rangle \in \mathcal{H}, \exists \theta \in \mathbb{R}, |g\rangle \in G, |b\rangle \in G^\perp, |\psi\rangle = \sin(\theta) |a\rangle + \cos(\theta) |b\rangle$. If we then apply $Q = -I_{|\psi\rangle} I_G$ where we assume that I_G can efficiently implemented then

Theorem 3. The **Amplitude Amplification Theorem** states that Q is the rotation

$$\begin{pmatrix} \cos(2\theta) & -\sin(2\theta) \\ \sin(2\theta) & \cos(2\theta) \end{pmatrix}$$

in the $|g\rangle, |b\rangle$ basis.

Proof. algebra... □

Using this we can greatly improve the probability of seeing $|g\rangle$. In particular we need $n \approx \frac{\pi}{4\theta}$ rotations to get $|g\rangle$ with a probability $\mathbb{P}(|g\rangle) \geq \cos^2(\theta) \geq 1 - O(\theta^2)$. We also note that in order to implement G it is sufficient to span G by the computational basis and to have an efficiently computable indicator function of the basis for G . We can also implement $I_{|\psi\rangle}$ in $O(n)$ time.

Some remarks on amplitude amplification are that it is useful to prepare a state one ones (start with the uniform state and an indicator function of the desired state, and proceed). Also, if $\sin(\theta)$ is known exactly, then the algorithm can be modified, at little cost, to become deterministic, which is somewhat interesting [End of lecture 8].

3.3 Applications of Amplitude Amplification

We can use Amplitude Amplification to generalise Grover's algorithm to a search with multiple possible solutions, instead of just a unique solution. If there are k possible solutions then using G as the space spanned by this solution we get $\sin(\theta)\sqrt{\frac{k}{N}}$ and then we can proceed as usual.

A far more peculiar application, is that as long as the solution can be checked efficiently, we can use Amplitude Amplification to get a speedup on any general algorithm. If we assume our algorithm is implemented efficiently (poly time) as the quantum gate A and

$$A|0\rangle = \alpha|a\rangle + \beta|b\rangle$$

for $\alpha = \sin(\theta)$ then applying the usual approach to reduce the error rate we find that amplitude amplification is quadratically more efficient than just repeating A and measuring. In particular, we set $|\psi\rangle = A|0\rangle$. Furthermore, if the error probability is known, then we can - as before - modify the algorithm to be deterministic. This is peculiar fact, since it is as of yet unknown in classical complexity theory whether or not $P = BPP$...

A final application is **quantum counting** which is to count the size (dimension) of G . This combines Amplitude Amplification and Phase Estimation. Here we compute our usual $Q = -I_{|\psi\rangle}I_G$ where $|\psi\rangle$ is just the uniform state, and then we notice that the eigenvectors of Q are

$$|e_{\pm}\rangle = \frac{1}{\sqrt{2}}(|b\rangle \pm i|g\rangle)$$

for eigenvalues $\lambda_{\pm} = e^{\pm 2i\theta}$. We then observe that

$$|\psi\rangle = \sin(\theta)|g\rangle + \cos(\theta)|b\rangle = \frac{1}{\sqrt{2}}(e^{i\theta}|e_{+}\rangle + e^{-i\theta}|e_{-}\rangle)$$

which is an equally weighted $|e_{\pm}\rangle$ superposition. We can then use phase estimation to estimate $\phi_{+} = \theta/\pi, \phi_{-} = 1 - \theta/\pi$ but we know that $\theta/\pi = \frac{1}{\pi}\sqrt{|G|/N}$ which is generally small (certainly less than 1/2), and so we can determine θ

so $|G|$ with certainty. Well, not quite, if we use m qubit lines we can get an m -bit approximation to $\sqrt{\frac{k}{m}}$ with $2^m C - Q$ gates so 2^m queries to f . If we then get an additive error of $2^{-m} = \delta/\sqrt{N}$ in \sqrt{k} leaving us with k to error $O(\delta\sqrt{k})$ in $O(\sqrt{N}/\delta)$ queries. Classically, on the other hand, $O(N/\delta^2)$ queries are necessary (sample randomly and use bounds on the law of large numbers to get this estimate) [End of lecture 9].

4 Hamiltonian Simulation

Now to move to a completely different topic, we're going to discuss Hamiltonian simulation, or how to simulate quantum Hamiltonians on a quantum computer, which is probably where many of the earliest applications of quantum computing will be. On a classical computer, fundamentally we need $O(2^n)$ operations to simulate n qubits. On a computer, we hope to do this in polynomial time.

In particular, we are seeking to emulate the Schrödinger equation, which means that given H we'd like to compute $U = e^{iHt}$ since

$$\partial_t |\psi\rangle = -iH |\psi\rangle \implies |\psi\rangle = e^{-iHt} |\psi\rangle$$

Note that this problem is fundamentally different than the problem of identifying the quantum ground state of a system, which is known to be a QMA complete problem. Now, inevitably some error will be involved in these calculations, so we will use the matrix operator norm to analyse this error. In particular we seek U such that $|e^{-iHt} - U| < \epsilon$ for some ϵ .

This cannot be solved for an arbitrary Hamiltonian, which would imply a very large class of problems that can be formulated in this way could be solved this way as well. In particular, we see that a class of particularly common can be well approximated.

Definition 2. Hamiltonian H is k -local if it can be written as

$$H = \sum_{j=1}^m H_j$$

where each H_j acts on at most q qubits. Note here that $m \leq \binom{n}{k} = O(n^k)$ is polynomial in n .

Examples of such Hamiltonians include the Ising model, and the more general Heisenberg model (uses a general inner product instead of the relatively limited Ising interaction). In Chemistry, also, covalent bonds are all like local Hamiltonians.

Now, the fundamental challenge to this problem is that

$$e^{A+B} \neq e^A e^B$$

Nevertheless, we can derive suitable approximations. Let's start with

Theorem 4. The **Solovay-Kitaev** theorem states that if U is a unitary operation on q qubits then given a universal set of quantum gates, U can be approximated to ϵ using $O(\log^c(1/\epsilon))$ gates where $c < 4$.

Proof. The proof is omitted here, but the Nielsen and Chuang book covers it. It really is a result on Lie algebras... \square

Lemma 1. Given unitary operators U_i, V_i such that $\forall i, |U_i - V_i| < \epsilon$ then

$$|U_m \dots U_1 - V_m \dots V_1| < m\epsilon$$

This is mathematically easy (ES 2), but it is quite a profound result, since classically error accumulates exponentially, but in the quantum case it only accumulates linearly... [End of lecture 10]

Proposition 1. The sum of commuting k -local Hamiltonians H_j can be approximated by a circuit of length $O(m \text{ poly}(\log(m/\epsilon)))$ for any universal set of quantum gates.

Proof. Approximate each $e^{-iH_j t}$ and multiply them, then by Solovay-Kitaev and error accumulation applied to an error ϵ/m for each individual gate we get the desired result. \square

In the non-commuting case, we need one additional result.

Lemma 2. The **Lie-Trotter** product formula is that for $|A|, |B| < K$ and $K < 1$

$$e^{-iA}e^{-iB} = e^{-i(A+B)} + O(K^2)$$

which can be proven by considering the Taylor series of the exponential function. This also generalises to the Suzuki-Trotter formulas which get more accurate.

Now, to do our non-commuting case assume that $K < 1/m$ then we see that

$$\begin{aligned} e^{-iH_1} \dots e^{-iH_m} &= (e^{-i(H_1+H_2)} + O(K^2))e^{-iH_3} \dots e^{-iH_m} \\ &= (e^{-i(H_1+H_2+H_3)} + O(2^2 K^2))e^{-iH_4} \dots e^{-iH_m} \\ &= e^{-\sum_j H_j} + Cm^3 K^2 \end{aligned}$$

Now this quickly gets out of hand, accuracy-wise, so in order to improve our accuracy a bit as well, we consider doing N steps of equal size t/N instead of a one out calculation. So we use $H_j t/N$ instead and repeat. If this introduces error $\tilde{K} = \frac{Kt}{N}$ instead then we see that

$$Cm^3 \tilde{K}^2 < \epsilon/N \implies N > \frac{Cm^3 k^2 t^2}{\epsilon}$$

as the minimum number of steps. Implementing the circuit such that

$$\left| \left(e^{-iH_1 t/N} \dots e^{-iH_m t/N} \right)^N - e^{-\sum H_j t} \right| < \epsilon$$

gives us a circuit size $O(m^4(Kt)^2/\epsilon)$. Also, $m = O(n^k)$ means that we really have a circuit size $O(n^{4k}(Kt)^2/\epsilon)$. Using a standard universal set of gates we then find that by Solovay-Kitaev we get an additional logarithmic multiplicative factor $O(\log^c(|\mathcal{C}|/\epsilon))$ where $|\mathcal{C}|$ is the circuit size from before and $c < 4$. This is a modest increase in circuit size.

Finally, one might be surprised that it takes t^2 operations to simulate t amount of time on a quantum computer, but in factor, by refining the Lie-Trotter formula, we can make this $t^{1+\delta}$ for δ arbitrarily small. We probably cannot make it sublinear in t . That would be quite remarkable [End of lecture 11].

[include standard quantum gates in course summary]

5 The Harrow Hassidim Lloyd (HHL) Algorithm

The Harrow Hassidim Lloyd (HHL) algorithm describes a way of solving systems of linear equations exponentially faster on a quantum system. However, it should be noted that by “solving” a system, we do not mean finding all the coefficients of the solution. Instead we simply mean that we find a state that is a normalised linear combination of all the right coefficients, which can be used to test specific properties, however, from which the full solution cannot be fully recovered without extra work.

In particular, we seek to solve

$$Ax = b, b, x \in \mathbb{C}^N$$

for large N in $\text{poly}(n)$ time for $n = \log(N)$. And we aim to compute the value of $x^\dagger M x$ completely, which allows us to say something about the solution. Given the ubiquity of linear systems, this would also have many many applications, especially in solving PDEs, but also data mining, etc. The best classical techniques in this area are $\text{poly}(N)$ so this represents an exponential speedup.

The three parameters of interest to which this system depends are

- the system size N
- the error tolerance ϵ
- the **condition number** $k = \left| \frac{\lambda_{\max}}{\lambda_{\min}} \right|$ where $\lambda_{\max}, \lambda_{\min}$ are the largest and smallest eigenvalues of A respectively (by absolute value). Intuitively, $k = \infty$ for a singular matrix, so one can think of k being a measure of “how” singular a matrix is. Lower is better here.

The HHL algorithm then can solve $Ax = b$ under the following conditions

1. We require A to be Hermitian, however, one should note that this does not really restrict the algorithm at all, since for a non-Hermitian matrix A one can solve the system

$$\begin{pmatrix} & A^\dagger \\ A & \end{pmatrix} \begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} 0 \\ b \end{pmatrix} \implies y = 0, x = A^{-1}b$$

Note that doing this on a system size $2N$ does not fundamentally change the complexity of the algorithm.

2. k is small, or in particular, A is **well-conditioned** meaning that k is bounded by $\text{poly}(n)$ for increasing N .
3. One can implement Hamiltonian simulation of e^{iAt_0} (need e^{iA}) in $\text{poly}(n, t_0)$. This is the only real restriction on A . It works for example when A is k -local.
4. b is normalised and $\sum b_i |i\rangle$ can be computed in $\text{poly}(n)$ (data can be “loaded” efficiently - loading A efficiently is implicitly covered in requiring its Hamiltonian simulation). Note if b is not normalised, then as long as $\sum b_i^2$ can be computed in $\text{poly}(n)$ so that we can rescale before and after we should be fine.
5. $x^\dagger Mx$ and the corresponding measurement is efficiently implementable on n qubits. We also require that M is Hermitian, however, note that again this is no real restriction, since by using

$$M = K + iL, K = \frac{M + M^\dagger}{2}, L = \frac{M - M^\dagger}{2i}$$

we can get around this by repeating the algorithm twice, which does not increase complexity.

A first remark will be on the general set of matrices for which Hamiltonian simulation can be implemented. Here we remark that while the literature is quite involved, we find the following definitions helpful.

Definition 3. A is a **row sparse matrix** if A has $\text{poly}(n)$ non-zero entries in each row.

Definition 4. A is a **row s -sparse matrix** if A has no more than s non-zero entries in each row.

Definition 5. A is a **row computable matrix** if A is row s -sparse and \exists a $O(s)$ -time computation $C(i, k) = (j, A_{ij})$ which computes given $1 \leq i \leq N$ and $1 \leq k \leq s$ gives the k th non-zero entry of row i in matrix A and j the column number of that non-zero entry.

The following theorem was then shown in 2007 by Berry, Ahokas, Cleve and Sanders (0508139 on arXiv).

Theorem 5. The **Hamiltonian Simulation Property** states that a row s -sparse, row computable matrix A can have e^{iAt_0} computed to error ϵ in $O(ns^2t_0)$ time.

Note that in the statement of terms like s^α where α can be chosen to be arbitrarily small (e.g. due to Lie-Trotter-(Suzuki?) formulas, etc.) have been omitted.

Conveniently, this property is very common in PDE systems, and we also rarely require every state in a PDE system (perhaps we only need the final state), which is covered nicely here. [End of lecture 12]

The way in which HHL provides informatino about the system is that it allows us to calculate the quantity $M = x^\dagger M x$ for any Hermitian M to accuracy ϵ in $O(\log(N)s^2k^2/\epsilon)$ compared with $O(Ns\sqrt{k}\log(1/\epsilon))$ classically. You might note the exponential slowdown in $1/\epsilon$, however, subsequent improvements of HHL have improved this part of the algorithm to $\log(1/\epsilon)$. Now let's actually discuss how the HHL algorithm runs:

Let's assume that phase estimation, and Hamiltonian simulation can be implemented error free, and wlog assume $|\lambda_{max}| = 1$ and that k , the condition number, is known or upper bounded. Then let us assume that we have eigenvalue λ_i for eigenvector $|u_i\rangle$ meaning that if $|b\rangle = \sum b_i |i\rangle = \sum \beta_i |u_i\rangle$ so $|x\rangle = A^{-1} |b\rangle = \sum \frac{\beta_i}{\lambda_i} |u_i\rangle$. But A is Hermitian, not unitary, so cannot be implemented as a gate. Instead, we can implement $U = e^{iA}$ and then implement the $C - U$ gates via Hamiltonian simulation. We can then use a post-selection process to complete our "non-unitary gate."

As such our algorithm takes the form

1. Apply the phase estimation unitary estimator $U_{PE} |b\rangle |0\rangle = \sum \beta_j |u_j\rangle |\lambda_j\rangle$.
2. add a qubit, and then apply the controlled rotation

$$C - \text{rot} |\lambda_j\rangle |0\rangle = |\lambda_j\rangle (\cos(\theta_j) |0\rangle + \sin(\theta_j) |1\rangle)$$

where $\sin(\theta_j) = c/\lambda_j$, $\cos(\theta_j) = \sqrt{1 - c^2/\lambda_j^2}$ where $c \leq \min\{\lambda_j\}$ (e.g. $c = 1/k$) is some constant. This controlled rotation, it is shown on ES 3, can implemented in $\text{poly}(n)$ time. Finally, we measure this last register.

3. repeat this step until we measure a 1 (this is the "post-selection step"). Here we have the probability of getting a one as

$$\mathbb{P}(1) = P = \sum_j \left(\beta_j \frac{c}{\lambda_j} \right)^2 = \frac{1}{k^2} \sum_j \left| \frac{\beta_j}{\lambda_j} \right|^2 \geq \frac{1}{k^2}$$

since $\lambda_j \geq 1$ here. Note that this is influenced by the accuracy of our estimate of k . We consequently get the state

$$|V\rangle = \frac{1}{\sqrt{P}} c \sum_{j=0}^{N-1} \frac{\beta_j}{\lambda_j} |u_j\rangle |\lambda_j\rangle$$

Complexity-wise we find that if this is achieved with probability $1 - \eta$ then it takes $\log(1/\eta)/P = O(k^2)$ attempts, although after amplitude amplification, this becomes $O(k)$ times.

4. Now run U_{PE} for phase estimation in reverse and discard the $|\lambda_j\rangle$ register to get

$$|\hat{x}\rangle = \frac{c}{\sqrt{P}} \sum_{j=0}^{N-1} \frac{\beta_j}{\lambda_j} |u_j\rangle = \frac{c}{\sqrt{P}} |x\rangle$$

5. To estimate

$$\frac{c^2}{P} \langle x | M | x \rangle = \frac{c^2}{P} \mu$$

we just perform many repeated measurements of the observable M in this state, which can be somehow implemented? By the Chernoff-Hoeffding bound, $O(\log(1/\eta)/\xi^2)$ measurements suffice to estimate μ to accuracy ξ with probability $1 - \eta$. From our system, we know the value of c , but how do we get the value of P ? This is done quite crudely, by noting that $P = 0 \times P_0 + 1 \times P_1$ is the expected value of the ancilla measurement in the first post-selection step of HHL. If during this post-selection step, we measure P as well, using the same Chernoff-Hoeffding bound we get the same result. [End of lecture 13]

5.1 The Accuracy of the HHL Algorithm

[this lecture - lecture 14 - on the accuracy bounds of HHL is non-examinable]
First of all, we remark that a more detailed discussion of all the error bounds can be found in the original HHL paper. For simplicity, we assume here that $|b\rangle$ can be found exactly, and that the controlled rotation we use can be implemented exactly. Then we note that phase estimation is the most troublesome part of the whole algorithm. In particular, if we use n qubit lines, and we assume that we get λ' the closest n -bit approximation to λ (so up to error $\eta = 2^{-n}$) then we see that for $U = e^{iA}$, $C - U = e^{i\tilde{A}}$ for

$$\tilde{A} = \begin{pmatrix} 0 & 0 \\ 0 & A \end{pmatrix}$$

(where each 0 is an $n \times n$ matrix) so we just need to be able to calculate $e^{i\tilde{A}t}$ for $t = 1, \dots, 2^{n-1}$, $t_0 = 1 + \dots + 2^{n-1} = 1/\eta$. The result we get from our algorithm is

$$|\hat{x}'\rangle = \frac{1}{D'} \sum \beta_j \frac{c}{\lambda'_j} |u_j\rangle, D' = \sqrt{P'} = \sqrt{\sum_j |\beta_j|^2 \frac{c^2}{\lambda_j'^2}}$$

Now, here we note that

1. If λ has additive error, $\delta(\lambda) = \eta$ then $\delta(\frac{1}{\lambda}) \sim \frac{1}{\lambda^2} \delta(\lambda)$, and the associated relative error is (by using derivatives)

$$\frac{\delta(1/\lambda)}{1/\lambda} = \delta(\lambda)/\lambda = \eta/\lambda$$

2. Consequently if A', B' approximate A, B , respectively to relative error ξ , then $\frac{A'}{B'}$ approximates $\frac{A}{B}$ to relative error ξ as well.

This means that the relative error in $\frac{1}{\lambda_j}$ is $\delta(1/\lambda)/(1/\lambda) = O(\eta/\lambda) = O(k\eta)$. [actually k should be a κ]. Using $A' = \beta_j c/\lambda_j, B' = D'$ we see that the amplitudes of $|\hat{x}\rangle$ have relative error $O(k\eta) = \epsilon$. Consequently, $k\eta = \epsilon \implies \eta = \epsilon/k \implies t_0 = k/\epsilon$. Phase estimation then has complexity $O(\log(N)t_0 s^2) = O(\log(N)(k/\epsilon)s^2)$, which repeated with implitude amplification yields $O(\log(N)k^2 s^2/\epsilon)$ as required.

This analysis simplifies a lot. We assume that λ' is the closest we can get to λ , and of course, we assumed controlled rotation and $|b\rangle$ can be implemented error free. Nevertheless, this is the core of the result. It also reveals how we can improve things. Namely, the way we go from $1/\epsilon$ to $\log(1/\epsilon)$ is by replacing phase estimation entirely with something else called a “linear combination of unitaries” (see ES 3, Q 5, or for the paper arXiv’s 1312.1414). We will also note that the original HHL paper also extends the algorithm for ill-conditioned matrices, which we do not discuss here, and improves performance somewhat by starting with a non-uniform superposition of states. [End of lecture 14]

6 Clifford Computations and the Classical Simulation of Quantum Computing

We’ve now wrapped up all the algorithms we will be covering, so let’s take a look at some more theoretical concepts. A big question remains the interest in trying to identify what property precisely allows quantum computing to outperform classical computing. It is still uncertain to what extent this is true. But some of the key concepts are the following. A **classical simulation of a quantum computation** is a collection of

- a quantum circuit C , perhaps described as a list of n qubit gates, with circuit size $N = \text{poly}(n)$
- a description of the input (as product states or computational basis states?)
- certain nominated output states

We then say that we have a **weak simulation** if we can sample the output distribution classically, and we have a **strong simulation** if we can calculate the P_0, P_1 (efficiently - so in polynomial “time”) classically. Certainly strong implies weak. (the multiline output definition of a strong simulation is slightly

more involved and mentioned on ES 3, but not in lectures). If this were always possible, then quantum computation would never give more than a polynomial improvement. Of course, quantum computing only ever achieves weak simulation, so demanding strong simulation is quite a tall order. One might note that one can always write down a direct strong simulation, but it may not be efficient. More interestingly, it can be shown that if every step in a computation can be written as a product state then we can always find an efficient direct strong simulation.

This result also has the nice conceptual interpretation that without entanglement, quantum computing has no real advantage. However, it is unclear whether quantum entanglement is the source of that advantage, or a result of it.

6.1 Clifford Computations

A useful subclass of computations are called Clifford Computations. Let's first note that we can define the **Pauli group** P_n to be the group generated by Pauli matrices $\{I, X, Y, Z\}$ acting on individual qubits in the n -qubit space together with multiplication by $\{\pm 1, \pm i\}$. This is certainly a finite subgroup of $U(2^n)$. Consequently we can define **Clifford Operations** C on n qubits as any unitary operator that conjugates Pauli operators to Pauli operators (ie. the normaliser of P_n in $U(2^n)$). These play an important theoretical role, but are also particularly important in quantum error correction. Clifford computations include Pauli matrices, the phase gate

$$S = \begin{pmatrix} 1 & \\ & i \end{pmatrix}$$

the controlled X or $C - X$ gate, the swap and Hadamard gate. In fact, one can show that

Theorem 6. The H, S and $C - X$ gates generate the Clifford computations.

(no proof given). A Clifford computation, in case it is not clear, is a circuit of Clifford gates/operations. The main result of this section then is the

Theorem 7. Gottesman-Knill theorem (variant) which states that a Clifford circuit C on n qubits of size $N = \text{poly}(n)$ (measured in number of H, S and $C - X$ gates) with input as any product state, and output a measurement on the first qubit line can always be strongly simulated.

There are many proofs, and the original inspiration arises from the so-called stabiliser formalism (see the Nielsen and Chuang book for details). [End of lecture 15]

Proof. Write $Z_1 = Z \otimes I \otimes \cdots \otimes I$ for $Z = |0\rangle\langle 0| - |1\rangle\langle 1|$, then $P_0 - P_1 = \langle \psi_f | Z_1 | \psi_f \rangle = \langle \psi_i | C^\dagger Z_1 C | \psi_i \rangle$ for initial and final states ψ_f, ψ_i . Now if as a sequence of H, S and $C - X$ gates $C = C_N \dots C_1$ then Z_1 as a difference of

Pauli operations when conjugated by each C_i becomes a different Pauli operator. Consequently, $P_0 - P_1 = \langle \psi_i | \tilde{P}_1 \otimes \cdots \otimes \tilde{P}_n | \psi_i \rangle$. Since the $H, S, C - X$ gates each act on no more than 2 qubit, this Pauli operation can be computed efficiently, leaving for $|\psi_i\rangle = \langle a_1 | \dots \langle a_n |$

$$P_0 - P_1 = \prod_{i=1}^n \langle a_i | \tilde{P}_i | a_i \rangle$$

so using $P_0 + P_1 = 1$ we get our result. \square

So in these Clifford circuits, quantum computers offer no advantage. Now let's extend our scope to include intermediate measurements in Clifford circuits. We then get the two scenarios

- a circuit with measurements is **adaptive** if the circuit can change based on measurement outcome.
- a circuit with measurements is **non-adaptive** if the circuit cannot change based on measurement outcome.

We then find that when we have product state inputs, and single-state outputs, we can strongly simulate non-adaptive circuits, whereas adaptive circuits achieve full quantum computation. We note that when we are restricted to the computational basis, then the adaptive case can be weakly simulated. For more see Jozsa and van den Nest (arXiv 1305.6190). An outline of the proof is as follows:

Proof. In the non-adaptive case, we can easily replace measurement with a fully unitary circuit by replacing a measurement on line k with a $C - X$ gate connected to a new line starting in the $|0\rangle$ state, with line k doing the controlling. The new line serves no other purpose, and is discarded after the copying. The effect of this is to at least for weak simulation to separate states, allowing us to do a single measurement at the end.

In the adaptive case, we use the Bravyi-Kitaev idea of so-called magic states. We use the fact that if we adjoin the non-Clifford T -gate to the set of Clifford operations

$$T = \begin{pmatrix} 1 & \\ & e^{i\pi/4} \end{pmatrix} = \sqrt{S}$$

we gain universal quantum computing. As such we start with a so-called “magic state” $|A\rangle = \frac{1}{\sqrt{2}}(|0\rangle + e^{i\pi/4}|i\rangle)$, then we can implement a T gate in an adaptive Clifford circuit using the so-called “ T -gadget” where we let any line k control a $C - X$ gate applied to a new ancilla line that starts in the $|A\rangle$ state. We then measure the outcome in the ancilla, and if we get a 1, we apply S to the k th line, and if not, we do not apply anything. For the purposes of weak simulation at least (?) this provides a T gate. \square

For some this result was somewhat disconcerting because it meant that the “fully classical” notion of measurement could be used to breach the gap between classical and quantum computing. I’m still not sure what I think about it...